

# Process Synchronization (CS-351)

# Agenda

- What is process synchronization?
- Producer Consumer Problem: Race Condition
- The critical region problem.
- Software Solutions: Semaphores, Peterson's solution
- Hardware Solutions: TestAndSet(), Swap()
- The Deadlock, Starvation, Priority inversion problems

# What is Process Synchronization?

- Cooperating processes can share data.
- Concurrent access to data may result in data inconsistency.
- **Process synchronization**: ensuring an orderly execution of cooperating processes that share data, in order to maintain data consistency.
  - Necessary to prevent race conditions.
- **Race condition**: when multiple processes (or threads) access and manipulate the same data concurrently and the outcome of the execution depends on the particular order in which the access takes place.
- Recall the bounded buffer implementation of the producer consumer problem (next slide).

# Race Conditions: Producer Consumer Problem

- **Bounded buffer** implementation (a wrap-around buffer):
  - Store the following variables in **shared memory**:

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    ...
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int counter = 0; //Keeps track of the # of available slots.
```

```
int in = 0; //First empty position in "buffer"
```

```
int out = 0; //First full position in "buffer"
```

```
//The buffer is empty when in == out
```

```
//The buffer is full when ((in+1) % BUFFER_SIZE) == out.
```

# Race Conditions: Producer Consumer Problem

## Producer code:

```
while (true)
{
    /* do nothing -- no free buffers */
    while (((in + 1) % BUFFER_SIZE == out);
        //Produce an item
        .....
        //Save the produced item
        buffer[in] = item;
        //Compute the next free index
        in = (in + 1) % BUFFER_SIZE;
        counter++;
}
```

## Consumer code:

```
while (true)
{
    //No items to consume
    while (in == out);
    // Consume an item
    item = buffer[out];
    // Compute the index of the next
    // item to consume
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    return item;
}
```

# Race Conditions: Producer Consumer Problem

- Machine code for updating the counter:

## Producer code:

$PI_1$ . register<sub>1</sub> = counter  
 $PI_2$ . register<sub>1</sub> = register<sub>1</sub> + 1  
 $PI_3$ . counter = register<sub>1</sub>

## Consumer code:

$CI_1$ . register<sub>2</sub> = counter  
 $CI_2$ . register<sub>2</sub> = register<sub>2</sub> - 1  
 $CI_3$ . counter = register<sub>2</sub>

- $PI$  = producer instruction
- $CI$  = consumer instruction
- The **concurrent** execution of the above codes, may correspond to many different **sequential executions**.
- The final value of the counter depends on the **order of execution**.

# Race Conditions: Producer Consumer Problem

## Producer code:

$PI_1$ . register<sub>1</sub> = counter  
 $PI_2$ . register<sub>1</sub> = register<sub>1</sub> + 1  
 $PI_3$ . counter = register<sub>1</sub>

## Consumer code:

$CI_1$ . register<sub>2</sub> = counter  
 $CI_2$ . register<sub>2</sub> = register<sub>2</sub> - 1  
 $CI_3$ . counter = register<sub>2</sub>

$PI$  = producer instruction

$CI$  = consumer instruction

- $PI_1$ : producer: register<sub>1</sub> = counter {register<sub>1</sub> = 5}
- $PI_2$ : producer: register<sub>1</sub> = register<sub>1</sub> + 1 {register<sub>1</sub> = 6}
- $CI_1$ : consumer: register<sub>2</sub> = counter {register<sub>2</sub> = 5}
- $CI_2$ : consumer: register<sub>2</sub> = register<sub>2</sub> - 1 {register<sub>2</sub> = 4}
- $PI_3$ : producer: counter = register<sub>1</sub> {counter = 6}
- $CI_3$ : consumer: counter = register<sub>2</sub> {counter = 4}
- What if the order of  $PI_3$  and  $CI_3$  is swapped (next slide)?

# Race Conditions: Producer Consumer Problem

## Producer code:

PI<sub>1</sub>. register<sub>1</sub> = counter  
PI<sub>2</sub>. register<sub>1</sub> = register<sub>1</sub> + 1  
PI<sub>3</sub>. counter = register<sub>1</sub>

## Consumer code:

CI<sub>1</sub>. register<sub>2</sub> = counter  
CI<sub>2</sub>. register<sub>2</sub> = register<sub>2</sub> - 1  
CI<sub>3</sub>. counter = register<sub>2</sub>

PI = producer instruction

CI = consumer instruction

- PI<sub>1</sub>: producer: register<sub>1</sub> = counter {register<sub>1</sub> = 5}
- PI<sub>2</sub>: producer: register<sub>1</sub> = register<sub>1</sub> + 1 {register<sub>1</sub> = 6}
- CI<sub>1</sub>: consumer: register<sub>2</sub> = counter {register<sub>2</sub> = 5}
- CI<sub>2</sub>: consumer: register<sub>2</sub> = register<sub>2</sub> - 1 {register<sub>2</sub> = 4}
- CI<sub>3</sub>: consumer: count = register<sub>2</sub> {counter = 4}
- PI<sub>3</sub>: producer: counter = register<sub>1</sub> {counter = 6}
- The value of counter is now 6 (4 in the previous slide)!



# The Critical Section Problem

- **Critical section:** a segment of process **code**, in which the process modifies **shared resources** e.g. updating a table or writing to a file.
- **No two processes** may execute in the **critical section at the same time** e.g. no two processes may update the same variable at the same time.
- **The critical section problem:** ensuring that only one process is executing in its critical section.

# The Critical Section Problem: Solution Requirements

- Requirements for solving the critical section problem:
  - **Mutual exclusion:** only **one** process may enter a critical section at a time.
  - **Progress:** no process executing **outside** of its critical section, may block other processes from entering theirs.
  - **Bounded waiting:** process **cannot be perpetually barred** by other processes, from entering its critical section.
- **Existing solutions:**
  - **Peterson's solution:** software-based solution.
  - **Synchronization hardware:** hardware-based solution.
  - **Semaphores:** specialized integer variables.

# Semaphores

- Are **easier** to use than hardware-based synchronization.
- An integer variable **S**, accessible only through **wait()** and **signal()** operations: for example:

```
wait(S)  
{  
    //Do nothing  
    while( S <= 0);  
  
    S--;  
}
```

```
signal(S)  
{  
    S++;  
}
```

- All modifications to **S** are **atomic**.

# Semaphores: Different Types

- Binary semaphore:

- The value can range only between 0 and 1.
- Also known as mutex (i.e. mutual exclusion) locks.
- The mutex is initialized to 1.
- Processes use the mutex as follows:

```
// Code of process P
do
{
    wait(mutex);

    // Critical section

    signal(mutex);

    // Remainder section
} while (TRUE);
```

```
wait(S)
{
    //Do nothing
    while( S <= 0);

    S--;
}
```

```
signal(S)
{
    S++;
}
```

## Student Participation: Disadvantages of Mutex ( Discussion)

- What are the possible issues of using mutex?

# Semaphores: Implementation

- The following implementation of mutex requires **busy waiting** i.e. constantly polling  $S$  in a **while** loop:

```
wait(S)
{
    //Do nothing
    while( S <= 0);

    S--;
}
```

- Fine if the critical section is **short**; otherwise it is **inefficient**.
- If there are **multiple instances** of the same resource, the other processes still need to wait.
- There are alternatives to busy waiting (next slide).

# Semaphores: Different Types

- **Counting Semaphore:**

- The value is **unrestricted**.
- Useful for coordinating access to a resource with **finite number of instances**.
- The semaphore is initialized to the maximum # **of instances**.
- When the process calls **wait()** it decrements the value of S:
  - If S remains positive, then the process may continue executing.
  - Otherwise, the process waits in a loop until S becomes positive.
- When the process calls **signal()**, it increments the value of S.

# Semaphores: Counting Semaphore Implementation

- With each semaphore associate a wait queue of processes (represented as linked list).
- Each entry in queue comprises:
  - **Value** (of type integer)
  - **Pointer** to **next** record in the list
- Redefine **wait()** and **signal()** functions as follows:

```
//Code of the semaphore  
typedef struct  
{ int value;  
  struct process *list;  
} semaphore;
```

```
wait(semaphore *S)  
{ S->value--;  
  if (S->value < 0)  
  { add this process to S->list;  
    block();  
  }  
}
```

```
signal(semaphore *S)  
{ S->value++;  
  if (S->value <= 0)  
  { remove a proc. P from S->list;  
    wakeup(P);  
  }  
}
```

- **wait()**: if **S->value** is non-positive: the calling process is added to the list and is **blocked** (switched to waiting state).
- **wakeup()**: if **S->value <= 0**, then a process is removed from the list and is **woken up** i.e. allowed to execute.



# Student Participation: Properties of the semaphores solution to the CS problem

- 1) The semaphore mutex can only have the values 0 or 1.
  - True
  - False
- 2) When each process runs on a separate CPU then mutual exclusion cannot be guaranteed.
  - True
  - False
- 3) If p1 runs on a faster CPU than the other processes then p1 can execute CS multiple times before another process gets a turn.
  - True
  - False

# Deadlocks

- Multiple processes waiting indefinitely for an event that can be triggered by only one of the waiting processes.
- **Example:** Let  $S$  and  $Q$  be two semaphores, set to the value of 1:

$P_0$	$P_1$
wait ( $Q$ );	wait ( $S$ );
wait ( $S$ );	wait ( $Q$ );
.	.
.	.
.	.
signal( $S$ )	signal( $Q$ )
signal( $Q$ )	signal( $S$ )

- What can possibly go wrong?

# Starvation

- A process is **never removed** from the semaphore queue in which it waits.
- May happen if the other processes keep grabbing the shared resource **before** the queued process.

# Priority Inversion


- When a **lower priority process** holds a lock and prevents a **higher priority** process from acquiring the lock (so it may access a critical region).
- **Example:** 3 processes **L**, **M**, **H** with priority:  $L < M < H$ .
  - **H** waits for a lock held by **L**, meanwhile process **M** becomes runnable (and does not require a lock).
  - Since **M** is the only non-blocked task, it will be scheduled to run and preempt **L** (**L** cannot finish its critical region unless it runs)! Hence, **H** has to wait until **L** is scheduled to run again, finishes the critical section, and releases the lock!
  - **Solution:** while **L** has holds lock, let it inherit the priority of **H** so it cannot be preempted. This is called **priority inheritance protocol**.

# Priority Inversion: Priority Inheritance

- If a **higher priority** process waits for a **lower priority** process to release the lock, the **lower priority process temporarily inherits** the priority of the **waiting process**.
- When the process **releases** the lock, its priority is **reverted** to the original.

# Priority Inversion

- **Example:** The Mars Pathfinder (1997) Mission:
  - The Sojourner rover would frequently hang because
    - A **high priority** task called bc\_dist was blocked by a **lower priority** task "ASI/MET", waiting to use a shared resource.
    - ASI/MET, in turn, was pre-empted by multiple **medium** priority tasks.
  - **Solution:** enable priority inheritance on all semaphores (change a global variable in the VxWorks OS).



# The Critical Section Problem: Peterson's Solution

- A software-based solution.
- Assumes two processes ( $P_i$  and  $P_j$ ).
- Assumes that the **LOAD** and **STORE** operations are **atomic** i.e. they cannot be interrupted.
  - **Not** a reasonable assumption on **modern** architectures.
- Two processes **share** the following variables:
  - **int** **turn**;
  - **bool** **flag**[2]
- **turn**: indicates **whose turn** it is to enter the critical section.
- **flag**: indicates if a process is **ready** to enter the critical section.
- **flag**[*i*] = true implies that process  $P_i$  is ready to enter its critical section!

# The Critical Section Problem: Peterson's Solution

- The two processes ( $P_i$  and  $P_j$ ) share the following variables:
  - **int turn**: whose turn it is to enter the critical section?
  - **bool flag[2]**: if a process is ready to enter the critical section?
  - **flag[i] = true** implies that process  $P_i$  is ready!

## Code of process $P_i$

```
do {  
    flag[i] = TRUE;  
    //j = i - 1  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = FALSE;  
        remainder section  
} while (TRUE);
```

## Code of process $P_j$

```
do {  
    flag[j] = TRUE;  
    //i = j - 1  
    turn = i;  
    while (flag[i] && turn == i);  
        critical section  
    flag[j] = FALSE;  
        remainder section  
} while (TRUE);
```



# The Critical Section Problem: Peterson's Solution

- **Observation:**

- Process  $P_i$  sets  $flag[i]$  to true, indicating that it's ready to enter its critical section.
- Process  $P_i$  sets  $turn$  to  $j$ , indicating that  $P_j$  may enter its critical section if it wishes to do so.

- **Result:**

- If both processes try to enter their critical sections at the same time, then  $turn$  will be set to both  $i$  and  $j$  at roughly the same time.
- The final value of  $turn$  will be determined by the process that gets to update  $turn$  last, thus preventing both processes from entering their critical section simultaneously.

# The Critical Section Problem: Synchronization Hardware

- **Solution for uniprocessor systems:** disable interrupts while modifying a shared variable - if no other instructions are being executed, then nothing can interfere with the variable modification operation.
  - Approach taken by nonpreemptive kernels.
  - Difficult on multiprocessor systems.
- **Special atomic hardware instructions:**
  - Must be atomic i.e. once started, always completes.
  - **TestAndSet():** can atomically test and then set the value of a variable.
  - **Swap():** can atomically swap the contents of two variables.

# The Critical Section Problem: Synchronization Hardware

- TestAndSet() instruction code:

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

- **Remember:** the whole function is guaranteed to be atomic i.e. non-interruptible.

# The Critical Section Problem: Synchronization Hardware

- Using TestAndSet() to provide synchronization:
  - Assume both processes **share** the variable **lock**:

Code of process  $P_i$

```
do {  
    // Do nothing - wait  
    while (TestAndSet (&lock ));  
  
    // Critical section  
    ...  
  
    lock = FALSE;  
  
    // Remainder section  
    ...  
} while (TRUE);
```

Code of process  $P_j$

```
do {  
    // Do nothing - wait  
    while (TestAndSet (&lock ));  
  
    // Critical section  
    ...  
  
    lock = FALSE;  
  
    // Remainder section  
    ...  
} while (TRUE);
```

# The Critical Section Problem: Synchronization Hardware

- **Swap()** instruction code:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

- **Remember:** the whole function is guaranteed to be atomic i.e. non-interruptible.

# The Critical Section Problem: Synchronization Hardware

- Using Swap() to provide synchronization:
  - Both processes share a **lock** variable (initialized to **false**).
  - Each process has a local variable **myWait**.

Code of process  $P_i$

```
do
{
    myWait = TRUE;

    while ( myWait == TRUE)
        Swap (&lock, &myWait );

    // Critical section
    lock = FALSE;

    // Remainder section
} while (TRUE);
```

Code of process  $P_j$

```
do
{
    myWait = TRUE;

    while ( myWait == TRUE)
        Swap (&lock, &myWait );

    // Critical section
    lock = FALSE;

    // Remainder section
} while (TRUE);
```