

Scheduling (CS-351)

Agenda

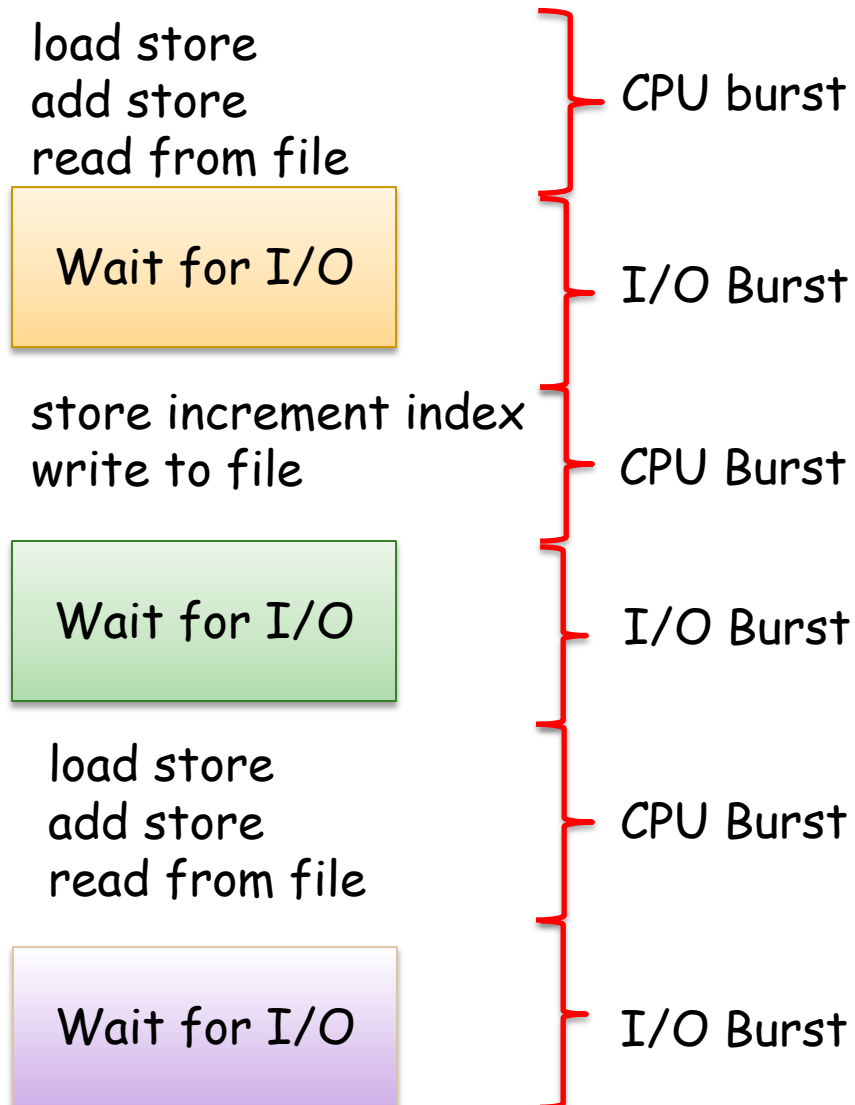
- Introduction to CPU scheduling.
- CPU Scheduling Algorithms: FCFS, SJF, SRT, RR, Priority queue, Multi-level queue
- Pthreads API

Basic Concepts: CPU Scheduling


- CPU is one of the **primary** computer resources.
- **CPU scheduling**: selecting the **next process** for execution on the CPU once the current process leaves the **CPU idle**.
- The goal of CPU scheduling is to **maximize** the degree of **multiprogramming** i.e., having some process running **at all times**.
- The success of CPU scheduling depends on an observed **property** of the processes:
 - CPU execution
 - I/O waiting

CPU-I/O Burst Cycle

- Processes alternate between **CPU bursts** (i.e. executing on the CPU) and **I/O bursts** (performing I/O).




Histogram of CPU-burst Times



- Large number of short CPU bursts and a small number of long CPU bursts.

The CPU Scheduler (a.k.a. The Short-Term Scheduler):

- Selects a process from the processes in memory that are ready to execute and allocates the CPU to the process.
- CPU scheduling decisions occur when:
 - 1. When a process switches from the running to the waiting state.
 - 2. When a process switches from the running to the ready state.
 - 3. When a process switches from the waiting state to the ready state.
 - 4. When a process terminates.
- Scheduling under 1 and 4 is nonpreemptive (e.g. Windows 3.11)
- All other scheduling is preemptive (e.g. Windows 95 and up and MAC OSX).



Preemptive vs Non-preemptive Scheduling

- **Preemptive scheduling:** process executing on the CPU can be interrupted in order to make way for another process.
- **Non-preemptive scheduling:** once a process gets the CPU, it runs to completion and cannot be interrupted.

Dispatcher

- A module that **gives control** of the CPU to the process selected by the scheduler.
 - Switches the **context**
 - Switches to **user mode**
 - **Jumps** to the proper location in the program to restart that program.
- Invoked during **every** process switch.
- **Dispatcher latency**: the time it takes for the dispatcher to **stop** one process and **start** another - should be **minimized**!

Scheduling Criteria

- Performance metrics of scheduling algorithms:
 - **CPU utilization**: keeping the CPU as busy as possible.
 - **Throughput**: the amount of processing that can be performed per unit of time.
 - **Waiting time**: the amount of time the process spends waiting in the **ready** queue.
 - **Response time**: amount of time it takes from when a request was **submitted** until the **first** response is produced.
 - **Running time**: the amount of **execution** time of the process
 - **Running time** = CPU burst time
 - **Turnaround time**: how long does it take to execute a process?
 - **Turnaround time** = waiting time + running time

Scheduling Algorithms: First-Come First-Served

- The process that **requests** the CPU first is **allocated** the CPU first.
- **Example:**

Process	Burst Time (milliseconds)
P ₁	24
P ₂	3
P ₃	3

- Suppose the processes arrive in the order: P₁, P₂, P₃:
 - **Gantt chart** (shows starting and ending times):



- **Waiting times:** 0 for P₁, 24 for P₂, and 27 for P₃.
- **Average waiting time:** $(0 + 24 + 27) / 3 = 17$ ms

Scheduling Algorithms: FCFS-cont.

- **Average Turnaround Time (ATT)** for the three processes using FIFO scheduling is 27.


	turnaround time	ATT(average turnaround time)
P1	$0+24=24$	$(24+27+30)/3 = 27$
P2	$24+3=27$	
P3	$27+3=30$	

Scheduling Algorithms: First-Come First-Served

- The process that **requests** the CPU first is **allocated** the CPU first.
- **Example:**

Process	Burst Time (milliseconds)
P ₁	24
P ₂	3
P ₃	3

- Suppose the processes arrive in the order: P₂, P₃, P₁:
 - **Gantt chart** (shows starting and ending times):



- **Waiting times:** 0 for P₂, 3 for P₃, and 6 for P₁.
- **Average waiting time:** $(6 + 0 + 3) / 3 = 3$ ms
- **ATT?**

Scheduling Algorithms: First-Come First-Served


- **Advantages:** simple to implement and understand.
- **Disadvantages:**
 - Average waiting time is often quite long.
 - **Convoy effect:** short process waiting behind a long process.
 - **Non-preemptive:** once the CPU is allocated to a process, that process keeps the CPU until it requests I/O or terminates.

Scheduling Algorithms: Shortest-Job-First

- Associates with each process the **length** of its next CPU burst. Uses these lengths to **schedule** the process with the shortest time. Can either be **preemptive** or **nonpreemptive**
- **Example (non-preemptive):**

Process	Burst Time (milliseconds)
P ₁	6
P ₂	8
P ₃	7
P ₄	3

- **Gantt chart:**



- **Waiting time:** 3 for P₁, 16 for P₂, 9 for P₃, 0 for P₄.
- **Average waiting time:** $(3 + 16 + 9 + 0) / 4 = 7\text{ms}$ (10.25 ms with FCFS).
- **ATT?**

Scheduling Algorithms: Shortest-Job-First

- Unlike nonpreemptive, will **interrupt** the currently executing process. Also known as **shortest-remaining-time-first (SRT)**.
- **Example (preemptive):**

Process	Arrival Time	Burst Time (millisecs)
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5



- **Waiting time:** $[(10-1) + (1-1) + (17-2) + (5-3)]/4 = 26/4 = 6.5$ ms
- **ATT?**

Scheduling Algorithms: Shortest-Job-First

- **Advantage:** is optimal - gives minimum average waiting time for a given set of processes.
- **Disadvantages:** we can only **estimate** the length of the next CPU request - a difficult task!
- One way to estimate the CPU burst of the process is to use **exponential averaging**.
- What if an **emergency** job but very long?

Student Participation: Determining the remaining time.


- The remaining time of a process depends on _____.
- 1) arrival time
 - True
 - False
- 2) total CPU time
 - True
 - False
- 3) real time in system
 - True
 - False
- 4) attained CPU time
 - True
 - False

Scheduling Algorithms: Priority Scheduling

- **Priority** is associated with each process. The CPU is allocated to the process with the **highest priority**. Can be **preemptive** or **nonpreemptive**.
- **Example** (smaller number = higher priority):

Processes	Burst Time (milliseconds)	Priority
P ₁	10	3
P ₂	1	1
P ₃	2	4
P ₄	1	5
P ₅	5	2

- All processes arrive at the **same time**.
- Gantt chart:




- **Average waiting time: 8.2 ms**
- **ATT?**

Scheduling Algorithms: Priority Scheduling

- **Preemptive priority** scheduling will preempt the CPU if the newly arrived process has a higher priority than the currently running process.
- **Starvation:** low-level priority processes may have to wait indefinitely.
 - **Example (rumor):** when MIT shutdown the IBM 7094 in 1973, they found a low-priority process that has been submitted in 1967 and had not yet run!
 - **Solution: aging:** gradually increase the priority of processes that wait in the system for a long time.

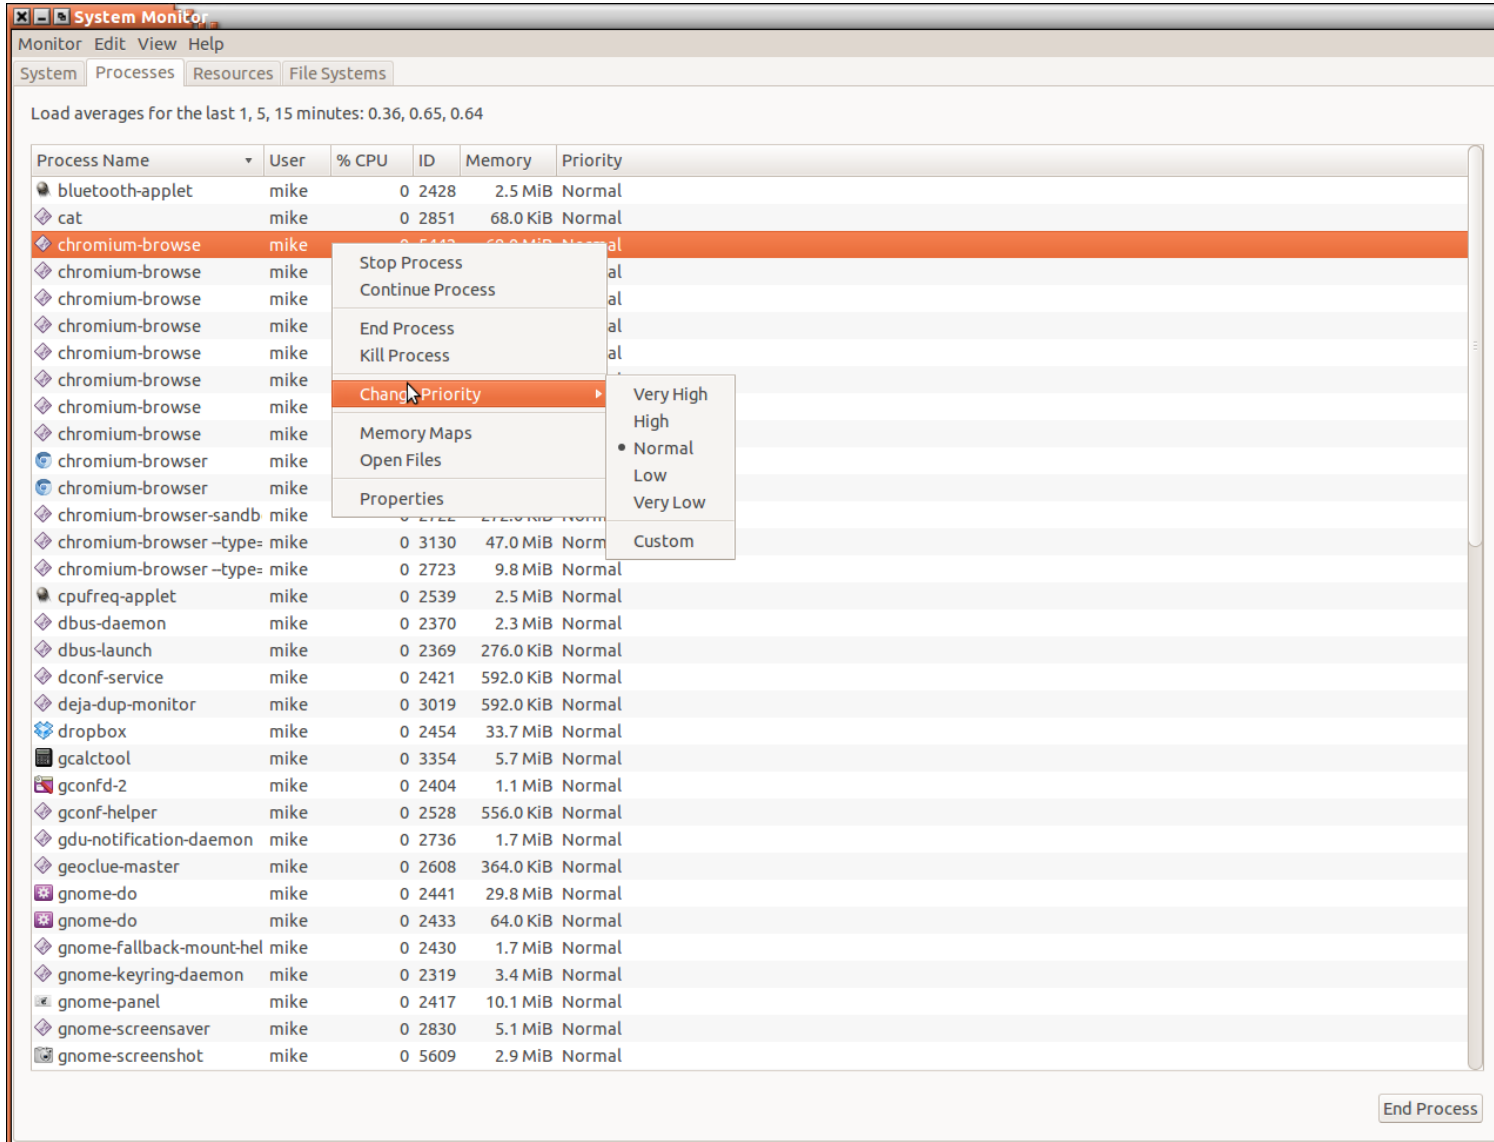
Scheduling Algorithms: Priority Scheduling

- Process priorities in Windows:



Scheduling Algorithms: Priority Scheduling

- Process priorities in Ubuntu Linux:



System Monitor

Monitor Edit View Help

System Processes Resources File Systems

Load averages for the last 1, 5, 15 minutes: 0.36, 0.65, 0.64

Process Name	User	% CPU	ID	Memory	Priority
bluetooth-applet	mike	0	2428	2.5 MiB	Normal
cat	mike	0	2851	68.0 KiB	Normal
chromium-browser	mike	0	5142	60.0 MiB	Normal
chromium-browser	mike	0	5143	60.0 MiB	Normal
chromium-browser	mike	0	5144	60.0 MiB	Normal
chromium-browser	mike	0	5145	60.0 MiB	Normal
chromium-browser	mike	0	5146	60.0 MiB	Normal
chromium-browser	mike	0	5147	60.0 MiB	Normal
chromium-browser	mike	0	5148	60.0 MiB	Normal
chromium-browser	mike	0	5149	60.0 MiB	Normal
chromium-browser	mike	0	5150	60.0 MiB	Normal
chromium-browser-sandbox	mike	0	5151	60.0 MiB	Normal
chromium-browser-type=	mike	0	3130	47.0 MiB	Normal
chromium-browser-type=	mike	0	2723	9.8 MiB	Normal
cpufreq-applet	mike	0	2539	2.5 MiB	Normal
dbus-daemon	mike	0	2370	2.3 MiB	Normal
dbus-launch	mike	0	2369	276.0 KiB	Normal
dconf-service	mike	0	2421	592.0 KiB	Normal
deja-dup-monitor	mike	0	3019	592.0 KiB	Normal
dropbox	mike	0	2454	33.7 MiB	Normal
gcalctool	mike	0	3354	5.7 MiB	Normal
gconfd-2	mike	0	2404	1.1 MiB	Normal
gconf-helper	mike	0	2528	556.0 KiB	Normal
gdu-notification-daemon	mike	0	2736	1.7 MiB	Normal
geoclue-master	mike	0	2608	364.0 KiB	Normal
gnome-do	mike	0	2441	29.8 MiB	Normal
gnome-do	mike	0	2433	64.0 KiB	Normal
gnome-fallback-mount-hel	mike	0	2430	1.7 MiB	Normal
gnome-keyring-daemon	mike	0	2319	3.4 MiB	Normal
gnome-panel	mike	0	2417	10.1 MiB	Normal
gnome-screensaver	mike	0	2830	5.1 MiB	Normal
gnome-screenshot	mike	0	5609	2.9 MiB	Normal

End Process

Scheduling Algorithms: Priority Scheduling

- Priorities in **Linux**: each process has a **nice** value specifying its priority:
 - Nice values range between -20 and 19.
 - **Lower** values indicate **higher** priority.
- Setting nice values:
 - **Example**: start program `ls` with nice value of 19:
 - `nice -n 19 ls`
 - **Example**: change nice value of already running process with id 1234, to 15.
 - `renice 15 -p 1234`

Scheduling Algorithms: Round-Robin (RR)


- Each process gets a small unit of CPU time i.e. **time quantum**. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- Given n processes and time quanta of q , Each process waits no longer than $(n - 1) * q$ time units.
- Performance is **sensitive** to the choice of the **time quantum**:
 - **Very large**: RR degenerates to **first-come first-serve**
 - **Very small**: q must be large in comparison to context switch latency, otherwise **overhead** is too high.

Scheduling Algorithms: Round-Robin (RR)

- Example (quantum = 4 millisecs):

Process	Burst Time (millisecs)
P ₁	24
P ₂	3
P ₃	3


- Gantt chart:



- **Waiting time:** P₁ waits for 6 (10-4), P₂ waits for 4, and P₃ waits for 7.
- **Average waiting time:** $17/3 = 5.66$
- **ATT?**


Scheduling Algorithms: Round-Robin (RR)

- How smaller quantum increases context switches:
 - Process time = 10



Scheduling Algorithms: Round-Robin (RR)

- How turnaround time varies with the time quantum:

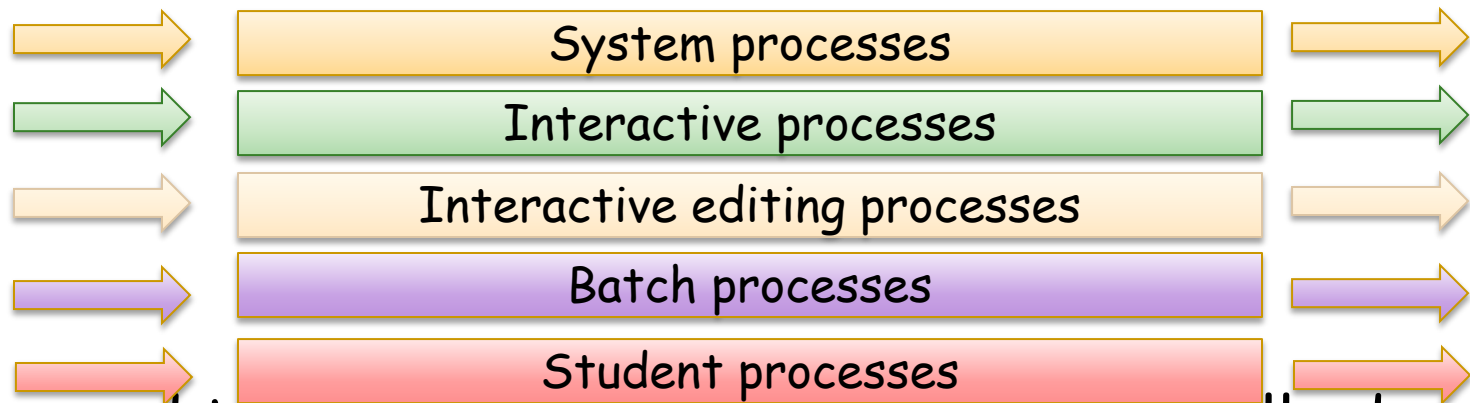


Process	Time
P ₁	6
P ₂	3
P ₃	1
P ₄	7

- Generally, the average turnaround time can be improved if most processes finish their **next** CPU burst in 1 quantum.
 - Example:** given 3 processes of 10 time units each and a quantum of 1, the avg. turnaround time is 29. Increasing the quantum to 10, reduces the avg. turnaround time to 20

Scheduling Algorithms: Multilevel Queue

- Partitions the ready queue into several separate queues. Processes are permanently assigned to queues, generally based on some property of the process e.g. memory size, priority, or type.
- Each queue has its own scheduling algorithm.
- Example:



- Also, need to schedule among the queues (usually done using fixed-priority preemptive scheduling algorithm).
 - I.e., from which queue should the next be selected?

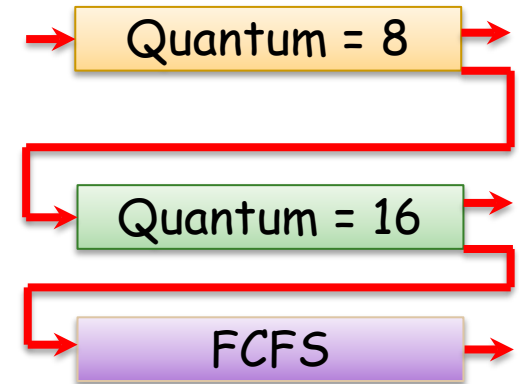
Scheduling Algorithms: Multilevel Feedback Queue

- Similar to multilevel queue scheduling, but the process can **move between queues**.
- The scheduler is defined by the following **parameters**:
 - **Number of queues**.
 - **Scheduling algorithms** for each queue.
 - Method used to determine when to **upgrade a process**.
 - Method used to determine when to **demote a process**.
 - Method used to determine which **queue** a process will enter when that process needs **service**.

Scheduling Algorithms: Multilevel Feedback Queue

- Example: 3 queues:

- Q_0 - RR with time quantum 8 milliseconds
- Q_1 - RR time quantum 16 milliseconds
- Q_2 - First-Come First-Served



- Scheduling:

- A new process enters queue Q_0 which uses round robin with quantum of 8.
- When it gets the CPU, the job receives 8 milliseconds of running time.
- If it does not finish in 8 milliseconds, the job is moved to queue Q_1 which uses round robin with quantum of 16.
- At Q_1 , the job receives additional 16 milliseconds.
- If it still does not complete, it is preempted and moved to queue Q_2 which uses FCFS scheduling.

Thread Scheduling: Pthreads

- User-level and kernel-level threads are scheduled differently.
- Many-to-one and many-to-many models, thread libraries schedule user-level threads to run on a **light weight process (LWP)** - a **virtual processor** on which the thread can be scheduled to run (or a kernel thread).
- **Process-contention scope (PCS): User** thread of a process competes for execution on a LWP (or kernel thread) with other user threads of the same process.
- **System-contention scope (SCS): Kernel** thread executing user threads of a particular process compete for execution on the physical CPU with other kernel threads executing user threads of the same process.

Thread Scheduling: Pthreads

- **Pthreads API** enables the developers to specify either PCS or SCS during thread creation:
 - **PTHREAD_SCOPE_PROCESS**: schedules threads using PCS scheduling i.e. each thread is bound to an available LWP.
 - **PTHREAD_SCOPE_SYSTEM**: schedules threads using SCS scheduling i.e. on many-to-many systems will create and bind an LWP for each user-level thread.

Thread Scheduling: Pthreads: Contention Scope

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    // Get the default attributes
    pthread_attr_init(&attr);
    // Set the scheduling algorithm to
    // PROCESS or SYSTEM
    pthread_attr_setscope(&attr,
        PTHREAD_SCOPE_SYSTEM);
    // Set the scheduling policy - FIFO, RT,
    // or OTHER
    pthread_attr_setschedpolicy(&attr,
        SCHED_OTHER);
    // Create the threads
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
}
```

```
// Now join on each thread
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i],
        NULL);
}
// Each thread will begin control
// in this function
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}
```


Thread Scheduling: Pthreads: Scheduling Policy

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5

int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* set the scheduling policy - FIFO, RT,
    or OTHER */
    pthread_attr_setschedpolicy(&attr,
    SCHED_OTHER);
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)

        pthread_create(&tid[i], &attr, runner, NULL);
}
```

```
/* now join on each thread */
for (i = 0; i < NUM_THREADS;
i++)
    pthread_join(tid[i], NULL);
}


/* Each thread will begin control in
this function */
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}
```

Multi-processor Scheduling

- CPU scheduling is **more complex** when multiple CPUs are available.
- Homogeneous processors within a multiprocessor.
- **Asymmetric multiprocessing:**
 - One processor **handles all scheduling, I/O processing**, etc. All other processors only execute **user code**.
 - Advantage: only one processor accesses the system data structures - don't need to worry about data sharing issues.
- **Symmetric multiprocessing (SMP):** each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes.
- **Processor affinity:** process has affinity for processor on which it is currently running.
 - **soft affinity:** a process **can** migrate between processors.
 - **hard affinity:** a process **cannot** migrate between processors

Multi-processor Scheduling

- Process affinities in Windows:



CPU Scheduling on NUMA Systems

- **Non-Uniform memory access systems (NUMA):** comprises combined CPU and memory boards. The CPUs on the board can access the memory on that board with **less latency** than on the other boards:
 - **Scheduling goal:** schedule a process on the CPU attached to the memory bank containing the process data.

