

Deadlocks (CS-351)

Agenda

- What is a deadlock?
- Resource-Allocation Graph
- Deadlock prevention.
- Deadlock avoidance.

What is a Deadlock?

- A set of processes $\{p_1 \dots p_n\}$ is **deadlocked** if each process in the set is **waiting** for an event that only **another** process in the set can cause.
- **Example:** A system with **one** printer and **one** DVD drive.
 - Process P_i is holding the DVD drive.
 - Process P_j is holding the printer.
 - If P_i requests the printer and P_j requests the DVD drive, a dead lock occurs.
- **Example:** a system has 3 CD RW drives D_1 , D_2 , and D_3 and processes P_1 , P_2 , and P_3 :
 - Each process P_i is holding drive D_i .
 - If each process requests access to another drive, a deadlock occurs.

What is a Deadlock?

- **Example:** "When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone" - law passed by the Kansas legislature in the early 20th century.
- **Example:** Let S and Q be two semaphores set to the value of 1, and P_0 and P_1 processes sharing these semaphores.

P_0	P_1
wait (Q);	wait (S);
wait (S);	wait (Q);
.	.
.	.
.	.
signal(S)	signal(Q)
signal(Q)	signal(S)

- The processes may become **deadlocked** after the second line.

What is a Deadlock?

- **Example:** Pthreads library uses `mutex` locks (which behave like `binary` semaphores) to provide mutual exclusion:
 - `pthread_mutex_t myMutex`: declares a variable called `myMutex` of type `mutex`.
 - `pthread_mutex_init(&myMutex, NULL)`: initializes `myMutex` and sets its state to "unlocked".
 - `pthread_mutex_lock(&myMutex)`: locks the mutex `myMutex`.
 - `pthread_mutex_unlock(&myMutex)`: unlocks the mutex `myMutex`.
 - If `myMutex` is already locked, then subsequent calls to `pthread_mutex_lock` will cause the calling thread to **block**.
 - If `myMutex` locking fails, then calling `pthread_mutex_lock` will return an error.

What is a Deadlock? - One Mutex example

```
#include <pthread.h>
#include <stdio.h>

/* This data is shared by the thread(s) */
int count = 0;           //The counter
                           variable.
pthread_t t1, t2;        //The thread variables.
pthread_mutex_t myMutex; //The mutex

/* the thread */
void *runner(void *param);
int sum = 0;

int main(int argc, char *argv[])
{
    pthread_attr_t attr; /* Set of thread
                           attributes */
    /* Get the default attributes */
    pthread_attr_init(&attr);

    /* Initialize the mutex */
    pthread_mutex_init(&myMutex, NULL);

    /* create the thread */
    pthread_create(&t1, &attr, runner, NULL);
    pthread_create(&t2, &attr, runner, NULL);
```

```
/* wait for the thread to exit */
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
} //End of main()

/* The thread will begin control in this
function */
void *runner(void *param)
{
    /* Lock the mutex to allow only one
thread */
    pthread_mutex_lock(&myMutex);
    for (int i = 0; i < 10; ++i) count += 1
    /* Unlock the mutex */
    pthread_mutex_unlock(&myMutex);

    pthread_exit(0);
}
```

Compare to Pthread multithreading sum example

What is a Deadlock? - Two mutex example

```
/* thread_one runs in this function */
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;
pthread_t t1,t2;
pthread_attr_t attr;

int main()
{
    /* Initialize the mutex locks */
    pthread_mutex_init(&first_mutex, NULL);
    pthread_mutex_init(&second_mutex, NULL);

    /* Get the default attributes */
    pthread_attr_init(&attr);

    /* create the thread */
    pthread_create(&t1, &attr, do_work_one, NULL);

    pthread_create(&t2, &attr, do_work_two, NULL);
    pthread_join(t1,NULL);
    pthread_join(t2,NULL)
}
```

Deadlock is possible if **thread_one** acquires **first_mutex** while thread two acquires **second_mutex**.

```
/* t1 runs in this function */
void *do_work_one(void* param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);

    /** Do some work **/

    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}

/* t2 runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);

    /** Do some work */

    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

Why do Deadlocks Occur?

- A deadlock can arise only if ALL of the following conditions are met:
 - **Mutual exclusion:** the resources involved are non-sharable i.e. if process P_1 requests resource R held by process P_2 , then P_1 must wait for P_2 to release the resource.
 - **Hold and wait:** a process must be holding at least one resource and waiting to acquire additional resources that are currently being held by other processes.
 - **No preemption:** resources cannot be preempted; a resource can be released only voluntarily by the process holding it, after that process has completed its task.
 - **Circular wait:** a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes must exist such that P_0 is waiting for a resource held by P_1 , P_1 is waiting for a resource held by P_2, \dots , P_{n-1} is waiting for a resource held by P_n , and P_n is waiting for a resource held by P_0 .

Resource-Allocation Graphs


- Used to precisely describe **deadlocks**.
- A set of **vertices** V and a set of **edges** E .
- V is **partitioned** into two types of vertices:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the **processes** in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all **resource** types in the system
- **Request edge**: directed edge $P_i \rightarrow R_j$.
- **Assignment edge**: directed edge $R_j \rightarrow P_i$.

Resource-Allocation Graph


- Process:




- Resource type with 4 instances:



- P_i requests an instance of R_j :




- P_i is holding an instance of R_j :




Resource-Allocation Graph - Example

- **Example:**
 - Process P_1 requests an instance of resource R_1 .
 - Process P_2 requests an instance of resource R_3 .
 - Instances of resource R_2 are assigned to processes P_1 and P_2 .
 - Instance of resource R_3 is assigned to P_3 .



Resource-Allocation Graph - Cycle


- **Example:**
 - P_3 now requests an instance of resource R_2 .




- The graph contains **2 cycles**:
 - $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$
 - $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$

Resource-Allocation Graph - Cycle vs Deadlock

- If graph contains **no cycles** then **no deadlock** is possible.
- If graph **contains** a cycle:
 - If only **one instance** per resource type, **then deadlock**.
 - If **several instances** per resource type, then **possibility** of deadlock.




Cycle: $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2 \rightarrow P_1$
and all resources in the cycle
have a **single** instance = **definite**
DEADLOCK!



No Cycles = definitely NO DEADLOCK!

Resource-Allocation Graph



Cycles:

$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$


$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$ and

Resource R₂ is in the cycle and has multiple instances = **possible deadlock!**

Actuality: P₁, P₂, and P₃ are **deadlocked**:

- P₂ is waiting for the R₃, held by process P₃.
- P₃ is waiting for either P₁ or P₂ to release R₂.
- In addition, P₁ is waiting for P₂ to release R₁.

Resource-Allocation Graph - Another example



Blocked, but not deadlocked

Cycle: $P_1 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

and R_1 and R_2 are in the cycle and have multiple instances = **possible deadlock!**

Actuality: No deadlock.

- Process P_4 may release its instance of R_2 .
- R_2 can then be allocated to P_3 , breaking the cycle.

Student Participation: Number of processes and resources needed for a deadlock

- Deadlock is possible when
 - A) the number of processes is greater than 1, regardless of the number of resources.
 - B) the number of resources is greater than 1, regardless of the number of processes.
 - C) the number of processes and the number of resources are both greater than 1.

Methods for Handling Deadlocks

- **Intuition:**

- Ensure that the system will **never** enter a deadlocked state.
- **OR allow** the system to enter a deadlocked, **detect** the deadlock, and **recover**.
- **OR pretend** that deadlocks **cannot happen** (the most popular approach i.e. used in Windows and Unix):
- Another words, it's up to application developers to write deadlock-free applications.

Methods for Handling Deadlocks

- **Specific Approaches:**

- **Deadlock Prevention:** making ensure that **at least one** of the four conditions necessary for the deadlock does not hold.
- **Deadlock avoidance:**
 - Require process provide information about resources it will need in the **future**.
 - Use this information to delegate resources in a way that will avoid deadlocks.
- **Deadlock detection and recovery:** allow deadlocks to happen, but have means of **recovering** from them.

Deadlock Prevention

- Must ensure that **at least one** of the following four conditions necessary for a deadlock does not hold.
- **1. Mutual exclusion:** only necessary for non-sharable resources:
 - **Example:** multiple processes cannot safely share a printer, but can a read-only file.
 - **Problem:** some resource (e.g. mutexes) are inherently non-sharable.

Deadlock Prevention

- Must ensure that **at least one** of the following four conditions necessary for a deadlock does not hold.
- **2. Hold and wait:** must guarantee that whenever a process requests a resource, it holds no other resources:
 - Require process to request and be allocated all its resources **before** it begins execution (e.g. do not allow system calls from the process until all resources are allocated).
 - **Problem:** low resource utilization, and may cause starvation.

Deadlock Prevention

- Must ensure that **at least one** of the following four conditions necessary for a deadlock does not hold.
- **3. No preemption:** if the process requests a resource that cannot be allocated immediately:
 - 1. **Preempt** all resources held by the process.
 - 2. **Add** the released resources to the list of resources requested by the process.
 - 3. **Restart** the process when all resources are **available**.
- **Problem:** only applicable to resources whose state can be **easily saved and restored**:
 - **Example:** CPU state, registers, etc. can be easily saved and restored. mutex/semaphore state cannot.

Deadlock Prevention

- Must ensure that **at least one** of the following four conditions necessary for a deadlock does not hold.
- **4. Circular wait:** impose a **total ordering** of all resource types, and require each process to request resources in an increasing order of enumeration.
 - Can be proved correct by contradiction.
 - **Example:** if two processes wants to use a tape drive and a printer, all the processes must first request a tape drive and then a printer (and not vice-versa).

Student Participation:

Eliminating the conditions for deadlock.

- 1) To guarantee that deadlock is impossible, the hold-and-wait condition must be eliminated for
 - at least 2 processes.
 - all processes.
- 2) To guarantee that deadlock is impossible, the circular-wait condition must be eliminated for
 - at least 1 process.
 - all processes.
- 3) To guarantee that deadlock is impossible
 - eliminating either of the two conditions is sufficient.
 - both conditions must be eliminated.

Deadlock Avoidance

- Next, we study an alternative technique called **deadlock avoidance**.